

D04.4 Initial prototype L4-based TC system

Project number	IST-027635		
Project acronym	Open_TC		
Project title	Open Trusted Computing		
Deliverable type	Report (see p 84/85 Annex 1 - Nature)		
Deliverable reference number	IST-027635/DD04.4/FINAL 1.0		
Deliverable title	D04.4: Initial prototype L4-based TC system		
WP contributing to the deliverable	WP4		
Due date	Feb 2007 - M16		
Actual submission date	April 2007		
Responsible Organisation	TUD		
Authors	TUD (Bernhard Kauer, Carsten Weinhold, Christelle Braun, Alexander Böttcher), HP (David Plaquin)		
Abstract	This report describes an initial prototype of an L4-based trusted operating-system layer. The prototype is based on a snapshot of the L4 software components that we provide in OpenTC and it includes a first implementation of the Basic Management Interface (BMI), facilities to access TPMs, and a trusted boot loader. This deliverable includes the full source code of the L4-based prototype as well as compiled binaries that have been integrated into a demonstrator in the form of a bootable CD image.		
Keywords	WP4, prototype, microkernel, L4, L4Linux, BMI, OSLO		
Dissemination level	Public		
Revision	FINAL 1.0		
Instrument	IP	Start date of the project	1 st November 2005
Thematic Priority	IST	Duration	42 months

If you need further information, please visit our website www.opentc.net or contact the coordinator:

Technikon Forschungs-und Planungsgesellschaft mbH
Richard-Wagner-Strasse 7, 9500 Villach, AUSTRIA
Tel. +43 4242 23355 -0
Fax. +43 4242 23355 -77
Email coordination@opentc.net

The information in this document is provided “as is”, and no guarantee or warranty is given that the information is fit for any particular purpose. The user thereof uses the information at its sole risk and liability.

Table of Contents

1	Introduction	3
2	The Software.....	3
2.1	L4/Fiasco Microkernel.....	3
2.2	Architectures based on L4/Fiasco.....	4
2.3	L4Env.....	4
2.4	L4Linux.....	4
2.5	BMI – The Basic Management Interface.....	4
2.6	STPM.....	5
2.7	IPCMon.....	5
2.8	OSLO.....	6
2.9	Summary of Work Done Within OpenTC.....	6
3	Demo.....	8
3.1	General Remarks.....	8
3.2	Booting IPCMon Test.....	8
3.3	Booting STPM Test.....	8
3.4	Booting BMI Test.....	9
3.5	Booting OSLO Test.....	9
3.6	Booting Integrated Demo of the Complete Prototype.....	9
3.6.1	L4Con and Run.....	9
3.6.2	STPM, BMI and IPCMon test.....	10
3.6.3	L4Linux and Networking.....	10
4	Development.....	10
5	Outline – further steps.....	11
6	List of Abbreviations	12
7	References.....	12

1 Introduction

The WP4 work package of the OpenTC project aims to build two alternative trusted operating system (OS) layers that provide isolation and protection mechanisms based on virtualization. Both layers will have an interoperable Basic Management Interface called BMI with common semantics for policies and configuration data in order to manage protection domains. In these protection domains, applications and OSes such as virtualized Linux instances are executed and protected from each other.

One of these OS layers is based on an L4 microkernel, the other one is based on the Xen VMM. This report describes the initial prototype of the L4-based trusted OS layer. It documents the basic functionality for the trusted OS layer that we developed within the OpenTC project, such as communication control and facilities to access TPMs as well as a rudimentary implementation of the BMI.

2 The Software

This deliverable is based on the February 2007 OpenTC snapshot of our **L4** software, which is publicly available at the OpenTC project website [\[3\]](#). The provided source tree contains the **L4/Fiasco** microkernel, the **L4Env**-core packages that provide services on top of this microkernel, and a paravirtualized Linux kernel called **L4Linux**. Additionally, we present in this deliverable the first prototype of the L4 BMI implementation, an L4Env service providing access to TPM interfaces called STPM, and our trusted boot loader Open Secure LOader (OSLO). In the next sections, we briefly describe the aforementioned software and give an overview of the components that were newly developed or modified for OpenTC. Where applicable, we refer to further documentation, before we present a demonstrator of the L4-based trusted OS layer in the next chapter.

2.1 L4/Fiasco Microkernel

L4 is a second-generation microkernel originally developed by Jochen Liedtke. The L4 philosophy is to implement only a minimal and essential set of functionality in kernel space in order to facilitate the construction of robust and tiny kernels. The L4 mechanisms provided to user level tasks are isolation of tasks through separation of address spaces, execution instances known as threads, and primitives for inter process communication (IPC) that allow threads to cooperate. In contrast to monolithic kernels such as Linux, all other required functionality including memory and IO management, device drivers, file systems, and complex protocol stacks (e.g., for networking) are implemented as services running as isolated L4 user-level tasks.

Nowadays, the name L4 also stands for a whole family of microkernels that implement various versions of the L4 application binary interface (ABI). One of these implementations is **Fiasco**, which is being developed and maintained at TU Dresden. Fiasco is sometimes referred to as **L4/Fiasco** to emphasize the relation to L4. TU Dresden is maintaining L4/Fiasco to improve stability and performance. Furthermore, there are ongoing efforts to port it to more computer architectures and it is extended with security features such as IPC monitoring, which supports the restriction of IPC-

based communication among user level applications.

2.2 Architectures based on L4/Fiasco

L4/Fiasco is real-time capable and supports parallel execution of real-time and time-sharing applications. This includes **L⁴Linux**, a paravirtualized version of Linux that is capable of executing on the L4 microkernel. It can run unmodified Linux programs. Furthermore, the **Nizza Secure-System Architecture** [15] is based on a microkernel -- in our implementation on L4/Fiasco -- and allows execution of security-critical applications in parallel to untrusted applications such as one or more L⁴Linux instances.

2.3 L4Env

L4Env comprises a set of user-level services and libraries that facilitate the use of L4/Fiasco and provide a basis for building an operating-system on top of L4. It provides memory management with dataspaces, management of tasks, threads and IO resources as well as services to start new applications. The development environment for L4Env includes DICE [14], an IDL compiler used to automatically generate stub code for remote function calls based on Interface Description Language (IDL) files. These IDL files specify the interfaces provided by L4Env services as a set of functions, which which applications can call remotely across address space boundaries.

Documentation of the L4Env is available at our L4Env webpage [5] and in the documentation directory (doc/) of the binary version of the OpenTC snapshot. The L4Env-core interface document [6] describes which functionality of the L4Env is publicly supported and which functions an application developer should not rely upon.

2.4 L⁴Linux

L⁴Linux [13] is a port of the Linux kernel to the L4 microkernel ABI and L4Env. As a paravirtualized Linux running on top of the microkernel, L⁴Linux is executed in ring 3 of the processor without privileges for accessing hardware IO ports or attaching to IRQs, and without the permission to perform privileged operations such as interrupt masking. Thus, multiple instances of L⁴Linux can run side-by-side with other L4 applications such as real-time or security components. Nevertheless, L4Env services can grant specific L4Linux instances the necessary rights to access certain hardware devices directly (e.g., a hard disk controller).

Because the paravirtualization only affects the inner functions of the kernel, L⁴Linux provides binary compatibility to existing x86 Linux applications and can therefore be used with any PC-based Linux distribution.

L⁴Linux is regularly updated to the latest Linux kernel version. Within this deliverable the version from the February 2007 OpenTC snapshot is used, which is 2.6.19.

2.5 BMI – The Basic Management Interface

The cbbmi package provides a first view on the BMI implementation for L4Env. It is not yet feature complete with regard to the current status of the BMI specification, but

mainly a snapshot of the current development. In this early stage, the BMI supports the following operations:

- start a protection domain (PD) such as L⁴Linux or other L4 applications via the L4 Loader or a guest L⁴Linux
- send a shutdown signal to an L⁴Linux PD
- kill a PD
- list all started PDs

The user manual of the cbbmi package describes the structure of the BMI implementation as well as the command line parameters of the bmic command-line application. It is available in the documentation directory of the cbbmi package [4].

2.6 STPM

STPM is a server that provides a /dev/tpm like interface for other L4 components. It reuses Linux-compatible device drivers for TPMs, mainly written by TUD, by using the Device Driver Environment (DDE) for L4. DDE provides a Linux kernel-like environment so that drivers originally written for Linux can be isolated in L4 tasks without the need of a complete Linux kernel. For a deeper explanation of the DDE ideas see [1,2].

The following drivers are stored in the contrib/ directory of the sources and compiled into the STPM server:

tpm/	The original Atmel TPM driver by IBM
inftpm/	A driver for v1.1 Infineon TPMs by TUD
oslo/tis.c	A driver for v1.2 TPMs of Atmel, Infineon, STM and Broadcom by TUD
tis_oslo/	Wrapper for the OSLO TIS driver with a Linux driver interface by TUD

The main difference between these drivers and those available in the 2.6.x Linux Kernels is that the former do not require ACPI support in order to detect the which particular TPM is installed. ACPI support is currently still missing in the L4 software stack. Instead, these drivers probe for supported TPMs.

STPM is reported to work with the following TPMs:

- Atmel v1.1 and v1.2
- Infineon v1.1 and v1.2
- STM v1.2
- Broadcom v1.2

STPM will be the basis for TPM support in the final BMI implementation.

2.7 IPCMon

IPCMon provides IPC control for L4/Fiasco at the level of L4 tasks. It consists of a kernel extension that adds support for task capabilities to Fiasco and a user-level server called ipcmmon that manages simple access policies and handles capability faults.

IPCMon was developed within the partner project EMSCB and will be used as the basis

to implement access control in the BMI.

2.8 OSLO

The Open Secure LOader (OSLO) is the first publicly available boot loader that leverages AMDs skinit instruction to create a dynamic root of trust as detailed in the TCG Specification. An important advantage over a static root of trust is that the size and complexity of the trusted computing base (TCB) can be reduced. As this instruction can be executed after the initial boot-up phase performed by the BIOS, it is possible to keep the BIOS and the actual boot loader (e.g, GRUB) out of the TCB. OSLO and all trusted L4 components are loaded using an untrusted boot loader that starts OSLO, which in turn puts the system into a trusted state using skinit and then passes control to the L4/Fiasco boot strapper.

OSLO was developed by TUD within OpenTC specifically to use the new hardware features of the AMD Athlon64 X2 processors. Due to its small size (approximately 1,000 lines of code resulting in a 4 kB of binary code), it should be feasible to review and verify the correctness of OSLO.

The latest version of OSLO is part of the contrib/ directory in the STPM package, however it can also be found at [7].

2.9 Summary of Work Done Within OpenTC

In the remainder of this chapter, we summarize the work that we did within WP4 of the OpenTC project. First of all, TUD invested considerable resources into maintenance of all the software components that we provide as part of our regular OpenTC snapshots. This work involved various of bug fixes and improvements of documentation.

Furthermore, the prototype of the L4-based trusted OS layer contains newly added components and modifications to L4/Fiasco, L4Env, and L⁴Linux that we did based on requirements and on feedback from other OpenTC partners:

Source directory	Description of new functionality and modifications
l4linux-2.6/arch/l4/	<ul style="list-style-type: none">● implementation of software suspend in L4Linux as required for future OpenTC demonstrators● improved support for multiple L4Linux instances running concurrently (ability to transfer FPU states, control assignment of ISA DMA memory)● several bug fixes and minor improvements
kernel/fiasco/	<ul style="list-style-type: none">● faster IPC path, more efficient use of kernel memory● support to transfer FPU states● support for KIP memory descriptors● several bug fixes
l4/pkg/ore/ l4linux-2.6/drivers/net/ l4ore.c	<ul style="list-style-type: none">● major functional improvements: support for loopback-only networking, multiple ORe instances providing isolated virtual LANs● several bug fixes

l4/pkg/cbbmi/	<ul style="list-style-type: none">● newly implemented
l4/pkg/libsigma0/ l4/pkg/sigma0/	<ul style="list-style-type: none">● MTRR support (cacheable memory-mapped I/O)● use of KIP memory descriptors for improved robustness
l4/pkg/stpm/	<ul style="list-style-type: none">● major functional improvements: OSLO-based TIS driver, Linux wrapper for TIS driver
l4/pkg/stpm/contrib/oslo/ /	<ul style="list-style-type: none">● newly implemented
l4/pkg/l4io/	<ul style="list-style-type: none">● support for mapping TPM TIS area
l4/pkg/ipreg/	<ul style="list-style-type: none">● newly implemented
l4/doc/dev-overview/	<ul style="list-style-type: none">● new document on how to get started with L4 development

3 Demo

3.1 General Remarks

The provided bootable CD image can be run on real hardware or within a virtualization software such as QEmu [8] or VMware [9]. Booting is done with GRUB 0.97 [10,11]. The serial port can be used for debugging purposes, see [12] for details.

Note however that the delivered software does not have product quality, but is rather a proof of concept that shows the current status of the implementation as well as future directions of development.

In the following sections, we describe the different choices the user is presented in the GRUB boot loader's menu when booting from the demo CD.

3.2 Booting IPCMon Test

The IPCMon test shows that IPC in L4 can be restricted by using IPCMon. This test loads, beside the ipcmmon server, two applications: ipctest1 and ipctest2. The former periodically outputs log messages. The latter revokes the rights of ipctest1 to communicate with the log server after a few seconds. The log output after revocation is something like:

```
ipcmmon | ipcmmon_pagefault(): ipc D -> 7 DENIED!  
*ipctest1 | main(): I'm still printing stuff to LOG.
```

In the first line IPCMon reports that it denied the communication from ipctest1 to the log server, which has task ID 7. The star at the beginning of the second line indicates that the log output is not printed via the now unreachable log server, but the kernel debugger instead.

3.3 Booting STPM Test

The STPM test demonstrates that a TPM can be accessed by L4 programs. It runs the STPM server together with a program called tpmdemo. The interesting lines of the output look similar to those below:

```
stpm | tis_oslo: Atmel rev: 2  
tpmdemo | TPM version 1.1.0.0  
tpmdemo | 24 PCR registers are available  
tpmdemo | PCR-00: 99 37 C4 68 E2 12 7D 16 BF CB D8 14 AC FC F0
```

The first line shows that the tis_oslo driver has found an Atmel TPM. After that, the number of the PCRs is shown and their values are dumped. Other output, such as the key handles or information about the endorsement key is also displayed.

3.4 Booting BMI Test

The BMI test starts the same applications as before, but uses the cbbmi server and bmic tool to start them. It tests the ability to start new applications with the BMI interface.

```
bmic | (bmilib) get_data(stpm$)
cbbmi | bmi_setImage_component(): (bmi)#####data:(stpm$)
cbbmi | startPD_intern(): (bmi) INIT DATANEW(stpm)
```

The first line shows the command line to start with the bmic tool “stpm”, the second line shows that the cbbmi server received the BMI setImage() call and the last line is from the startPD() function of the BMI.

3.5 Booting OSLO Test

This test just boots OSLO alone. It can be used to check whether OSLO supports a particular machine. The interesting lines on a machine not being an SVM platform are:

```
OSLO: No SVM platform
PAMPLONA: no ext cpuid
PAMPLONA: no module to start
```

The last line is always present in this test, as there is nothing started afterwards.

3.6 Booting Integrated Demo of the Complete Prototype

The full demo resembles all the aforementioned single tests into an integrated scenario.

3.6.1 L4Con and Run

The demo is booted into L4con, a graphical console that can be used to multiplex the framebuffer output of multiple applications. Please note that L4con is not supported in OpenTC and used here only to ease the visualization of various running L4Linux instances.

Two applications are running: logcon and run. The former is showing the log output, also visible on the serial line. The latter can be used to start new applications. Switching between the L4con consoles is done by using the shift and the function_keys e.g. Shift+F1 to switch to logcon.

The help of run can be invoked by typing the single key h. New applications can be loaded by typing l. An L4Linux can be started by loading lx.conf.

3.6.2 STPM, BMI and IPCMon test

The STPM test can be performed here by executing:

```
bmic --start stpm
bmic --start tpmdemo
```

The output is shown on the serial port or on logcon. The bmic tool can be used either via RUN or within an L4Linux shell.

The IPCMon test runs similarly by executing:

```
bmic --start ipcmmon
bmic --start ipcmmon_test.cfg
```

3.6.3 L4Linux and Networking

Multiple L4Linux instances could be started via

```
bmic --start-linux vmlinux 15
```

The last parameter gives the amount of memory they can use in megabytes. The one before the name of the kernel.

The started L4Linux instances are connected via our virtual network switch (ORe). A broadcast ping such as

```
ping 192.168.0.255
```

can be used to show which L4Linuxes are started.

It is also possible to log into these virtual machines via a simple remote shell on port 22. Use for instance

```
echo cat /proc/cmdline | nc 192.168.0.11 22
```

to get the commandline of a remote kernel.

There is currently no easy policy enforcement possible to restrict that a L4Linux instance can connect to ORe. A device abstraction layer in the BMI could be implemented to simplify the management of the virtual devices. This part of the BMI could use IPCMon as underlying service.

4 Development

For developing software in the L4 environment the development overview [12] can be used as a starting point. The document briefly describes and gives pointers to the needed tools, the test environment and the way to write new components or to build existing ones.

5 **Outline – further steps**

The next steps are to complete the BMI implementation, so that it meets the final version of the BMI specification, which is still under development. We will further test and stabilize our software. Other possible orientations for future research are the field of dynamic root of trust and new hardware features such as IO-MMUs.

6 List of Abbreviations

TPM	Trusted Platform Module
BMI	Basic Management Interface
DICE	IDL Compiler
IPC	Inter Process Communication
L4	A second generation microkernel interface
Fiasco	An implementation of L4 by TUD
PD	Protection domain
L4 ABI	L4 application binary interface

7 References

- [1] Christian Helmuth: "[Generische Portierung von Linux-Gerätetreibern auf die DROPS-Architektur](http://os.inf.tu-dresden.de/papers_ps/helmuth-diplom.pdf)" - http://os.inf.tu-dresden.de/papers_ps/helmuth-diplom.pdf
- [2] Thomas Friebe: "[Übertragung des Device-Driver-Environment-Ansatzes auf Subsysteme des BSD-Betriebssystemkerns](http://os.inf.tu-dresden.de/papers_ps/friebe-diplom.pdf)" - http://os.inf.tu-dresden.de/papers_ps/friebe-diplom.pdf
- [3] TUD OpenTC Snapshot <http://os.inf.tu-dresden.de/opentc/>
- [4] Basic Management User Manual [/src/l4/pkg/cbbmi/doc/bmimanual.pdf](http://os.inf.tu-dresden.de/src/l4/pkg/cbbmi/doc/bmimanual.pdf)
- [5] [L4Env – Documentation](http://os.inf.tu-dresden.de/l4env/docu.xml) <http://os.inf.tu-dresden.de/l4env/docu.xml>
- [6] [L4Env-core Interface](http://os.inf.tu-dresden.de/opentc/download/l4env-core.pdf) <http://os.inf.tu-dresden.de/opentc/download/l4env-core.pdf>
- [7] [OSLO – The Open Secure LOader](http://os.inf.tu-dresden.de/~kauer/oslo/) <http://os.inf.tu-dresden.de/~kauer/oslo/>
- [8] [Qemu](http://www.qemu.org/) <http://www.qemu.org/>
- [9] [Vmware](http://www.vmware.com) <http://www.vmware.com>
- [10] [GRUB - The GRand Unified Bootloader](ftp://alpha.gnu.org/gnu/grub/) <ftp://alpha.gnu.org/gnu/grub/>
- [11] [GRUB patches](http://os.inf.tu-dresden.de/~adam/grub/) <http://os.inf.tu-dresden.de/~adam/grub/>
- [12] [Developing with L4](http://os.inf.tu-dresden.de/l4env/doc/html/dev-overview/) <http://os.inf.tu-dresden.de/l4env/doc/html/dev-overview/>
- [13] [L4Linux](http://os.inf.tu-dresden.de/L4/LinuxOnL4/) <http://os.inf.tu-dresden.de/L4/LinuxOnL4/>
- [14] [DICE](http://os.inf.tu-dresden.de/dice/) <http://os.inf.tu-dresden.de/dice/>
- [15] [Nizza Secure-System Architecture](http://os.inf.tu-dresden.de/papers_ps/nizza.pdf) http://os.inf.tu-dresden.de/papers_ps/nizza.pdf