

D4.6 Initial prototype Xen-based TC system

Project number	IST-027635		
Project acronym	Open_TC		
Project title	Open Trusted Computing		
Deliverable type	Report (see p 84/85 Annex 1 - Nature)		
Deliverable reference number	IST-027635/D4.6/FINAL 1.0		
Deliverable title	Initial prototype Xen-based TC system		
WP contributing to the deliverable	WP4		
Due date	Feb 2007 - M16		
Actual submission date	April 2007		
Responsible Organisation	CUCL		
Authors	CUCL (Derek Murray) POLITO (Gianluca Ramunno)		
Abstract	In this report, an initial prototype Xen-based trusted computing system is described. This is a modified version of the Xen VMM and tool-stack, which takes into account some of the requirements of the Basic Management Interface (BMI). Specifically, our modifications to Xen have involved disaggregating the function of domain building from the management domain into a smaller domain isolated from the management domain. This modified version of Xen may be built using the Proof of Concept LiveCD, and the included development domain, though these do not form part of this deliverable.		
Keywords	WP4, Xen, Prototype, BMI, domB, disaggregation		
Dissemination level	Public		
Revision	FINAL 1.0		
Instrument	IP	Start date of the project	1 st November 2005
Thematic Priority	IST	Duration	42 months

If you need further information, please visit our website www.opentc.net or contact the coordinator:

Technikon Forschungs-und Planungsgesellschaft mbH
Richard-Wagner-Strasse 7, 9500 Villach, AUSTRIA
Tel. +43 4242 23355 -0
Fax. +43 4242 23355 -77
Email coordination@opentc.net

The information in this document is provided “as is”, and no guarantee or warranty is given that the information is fit for any particular purpose. The user thereof uses the information at its sole risk and liability.

Table of Contents

1	Introduction	4
1.1	Summary of work done.....	4
2	Features for Trusted Computing in Xen.....	5
2.1	Existing features.....	5
2.1.1	Hypervisor isolation.....	5
2.1.2	Mandatory Access Control (MAC).....	6
2.1.3	Virtual Trusted Platform Module (vTPM).....	6
2.2	Security-related features under development.....	7
2.2.1	Dom0 disaggregation.....	7
2.2.2	Xen Security Modules.....	7
2.2.3	IOMMU support.....	7
3	Interaction with the Proof of Concept.....	9
3.1	Description of the PET Proof of Concept prototype.....	9
3.2	Description of the development domain.....	9
3.3	Building the prototype using the development domain.....	9
4	Conclusion.....	11
5	List of Abbreviations	12

1 Introduction

The OpenTC project is building an open trusted computing architecture that is based on the virtualisation of computers. By harnessing virtualisation, it is possible to isolate trusted and untrusted code in separate virtual machines (VMs), and be assured that the trusted code is not affected by untrusted code running in any other VM.

Xen is a Virtual Machine Monitor (VMM) or *hypervisor*, that was originally developed at the University of Cambridge Computer Laboratory. It is capable of running several different operating systems in parallel. In the first two versions of Xen, a guest operating system had to be modified in order to run on the hypervisor, a process known as *paravirtualisation*. In the most recent major version of Xen, version 3, support for *full virtualisation* was added: new machine instructions available in the latest processors from AMD and Intel allow closed-source operating systems, such as Microsoft Windows, to be run in virtual machines. This was not possible before by simply using paravirtualisation, because the required modifications of the source code of such operating systems were not feasible.

In this report, an initial prototype of a Xen-based trusted computing system is described. It is a modified version of the Xen VMM and tool-stack that addresses some of the requirements of the Basic Management Interface (BMI). Our modifications to Xen have included disaggregating the function of domain building from the management domain into a smaller domain “domB”. This domain is isolated from the management domain. The modified version of Xen may be built using the Proof of Concept LiveCD, and the included development domain, though these do not form part of this deliverable. The remainder of the report is structured as follows. In Section 2, the relevant features of Xen for trusted computing are described. In Section 3, the Proof of Concept application and development domain are introduced, and their relation to the present work is outlined.

1.1 Summary of work done

In order to distinguish the OpenTC efforts from independent developments by the Xen community, we list the work that was done for OpenTC. Note that other work was also carried out in this area, both under the XenSE project and as part of the open-source community's commitment to improving the security rating of the Xen VMM.

- Modified Xen in order to allow the delegation of domain building privileges to domains other than dom0
- Created a new guest operating system (domB) that has the ability to build other guest operating systems, which involved:
 - Porting the Xen C-based control library, libxc, to Mini-OS
 - Creating a driver (gntdev) for mapping granted pages into user-space
- Modified the existing dom0 tool stack to enable the use of domB, which involved:
 - Adding a function (xc_domb_linux_build) to the dom0 instance of libxc, which acts as a proxy for the domain-build operation, and communicates with domB
 - Modifying the Python-based Xen control daemon to allow domB to be used instead of the standard builder, by specifying this in the domain configuration file

2 Features for Trusted Computing in Xen

In this section, the security-related features of Xen are described. In subsection 2.1, the features of the current Xen VMM are described. Subsection 2.2 describes TC-specific features, being developed in this work package and elsewhere, that will improve the security of Xen.

2.1 Existing features

2.1.1 Hypervisor isolation

A typical Xen system comprises one or more guest operating systems, running in their own dedicated *domains*. The most important feature of Xen for trusted computing is the isolation between domains that is provided by the hypervisor. This is a fundamental requirement for enforcement of any information flow policy. Every Xen system has a management domain, called *dom0*. *This domain is a privileged domain regarding access to the hypervisor API (or hypercalls)*. Other unprivileged domains containing a guest operating system are referred to as *domU*. The isolation properties of the hypervisor ensure that no *domU* may access the memory of another *domU* or of *dom0*. At present, the hypervisor makes no attempt to address covert channels.

For paravirtualised guest operating systems, this is typically achieved by making changes to the page tables indirectly through the Xen hypervisor API. When it is necessary to modify a page table, the guest operating system must make a *hypercall* to Xen, which invokes the appropriate code within the hypervisor and checks that the update is valid before committing the change to the relevant page table entry. An update is valid if the updated machine frame number is a member of the allocation for the domain; therefore it is not possible to map arbitrary memory from other guests, thus preventing unauthorised sharing of data between two domains. When it is necessary to make a large number of updates, such as upon process creation, the individual updates may be batched into a single hypercall, for performance reasons.

For fully virtualised guest operating systems, it is not possible to modify the page table update code, so, instead, *shadow page tables* are maintained. The unmodified operating system makes updates to its page tables by writing to them directly as it normally does; the page fault mechanism then signifies to Xen that the page table has been modified, and the hypervisor records the change in its shadow copy. The difference between the operating system and shadow copies is that the shadow page tables are a mapping from virtual page numbers to machine frame numbers. The operating system page tables are a mapping from virtual page numbers to pseudo-physical frame numbers, which are allocated to give the guest operating system the appearance of a contiguous physical address space, beginning at address zero.

It is possible, however, for two guest operating systems to share memory in order to communicate. This is done in a controlled manner, through the *grant table* mechanism. A guest that wishes to share its memory with another guest can create an entry in its grant table, which details the machine frame to be shared and the domain with which to share it. This yields a *grant reference*, which the foreign guest may use in a hypercall to map the granted page. Therefore, the sharing of memory is controlled by the domain that owns the memory.

The other mechanism for communication between guest operating systems is through the use of *event channels*. An event channel is similar to an interrupt, and paravirtualised operating systems typically map events onto virtual interrupt requests. A common means of inter-domain communication in Xen is to build a ring buffer from a page of shared memory, and use an event channel to notify the foreign domain of pending items.

2.1.2 Mandatory Access Control (MAC)

The current version of Xen includes optional MAC functionality. This allows the administrator of a Xen host to put in place formal policies that dictate what resources may be shared.

As were described in the previous subsection, the two primitive mechanisms for sharing resources in Xen are grant tables (shared memory) and event channels. The advent of MAC in Xen has led to hooks being added in the routines that create grant table entries and event channels. When MAC is enabled, these hooks call into the Access Control Module, which queries the current policy and returns a decision on whether or not the sharing should be allowed. For efficiency, these decisions are cached until the next time the policy changes. This architecture is necessary for the support of the WP5 security services implementation by allowing enforcement of the security policies at the hypervisor level while letting the policies being managed by the critical security services of the whole system.

The current implementation of MAC includes two example policies. The first is a Chinese Wall Policy, which is used to ensure that a guest cannot communicate with guests in two or more “conflict of interest classes”. The second is Simple Type Enforcement, which allows the administrator to assign a type to each guest, and enforce sharing rules between types.

2.1.3 Virtual Trusted Platform Module (vTPM)

Xen currently provides a virtual TPM device to guest operating systems, using the same split-driver model that is used for network and block devices. A front-end driver runs in domU, and provides a similar interface to a hardware TPM as defined by the Trusted Computing Group (TCG). The front-end driver passes requests through to the back-end driver that runs in dom0. A daemon process, running in dom0, creates vTPM instances for each domU when necessary. There is a one to one mapping between a domU and a specific vTPM, thus each domU has the impression it has access to a traditional hardware TPM.

Although this architecture provides TPM functionalities to each domU, there is no verifiable trust relationship between each vTPM and the hardware TPM of the platform. Furthermore, the trusted computing base for the current vTPM implementation comprises the hypervisor and dom0 as a whole, which makes it more difficult to trust vTPM implementations. These issues will be addressed in future versions of the prototype once the TPM Virtualisation architecture and its integration with the BMI is finalized in forthcoming WP4 deliverables.

2.2 Security-related features under development

2.2.1 Dom0 disaggregation

In its current release version, the implementation of Xen relies on the ability of the management domain (dom0) to map arbitrary pages of physical memory in order to manage the life-cycle of the other domains. This is used by the domain builder, XenStore and the console applet, present in dom0. This is an undesirable situation as it increases significantly the amount of the software that needs to be trusted in order to ensure isolation between the guest domains. Thus, the totality of the code of dom0 has to be trusted by a third party to have the assurance the isolation policies between the guest domains are actually being enforced. We therefore aim at providing Xen with a finer-grained privilege model, which will be used to remove this privilege from dom0.

In order to support domain building, a new paravirtualised guest with the ability to map arbitrary pages (*domB*) has been created. The code for building a domain has been ported to this guest, which is based on *Mini-OS*, a minimal operating system that is included with Xen, and which forms a much smaller trusted computing base than dom0. To support communication with XenStore and the console device, the grant table mechanism will be modified to include pre-defined grants to the pages that are used for XenStore and the console driver. In addition, it will be necessary to perform operations such as domain save and restore indirectly through a privileged domain, such as domB.

A prototype version of domB, along with the necessary modifications to Xen, has been included on the accompanying CD. **N.B.** This is prototype software, and should not be used in a production system.

This new DomB architecture will be used as the basis for the implementation of the Basic Management Interface (BMI) service, which is also being developed as part of WP4 in OpenTC. With this architecture, the BMI service will benefit from a small TCB and therefore provide local software and third parties with a higher level of trust in its behavior and the properties of the whole system.

2.2.2 Xen Security Modules

A future version of Xen will replace the current MAC code with the Xen Security Modules (XSM) framework. The XSM infrastructure has two main facets. The first is a hook infrastructure, in which hooks have been added to all potentially security-critical operations in the Xen VMM. The second is the security module itself, which installs functions for each of the hooks provided by the infrastructure. It is possible to specify the module to be used at boot time, as an image loaded by the bootloader.

XSM supersedes the current MAC implementation, by allowing more-flexible policies to be specified. The proposed implementation of XSM includes two pre-defined modules: one emulates the current MAC implementation, and the other implements the Flask security policy language, which is employed by SELinux.

2.2.3 IOMMU support

Once the ability of dom0 to map arbitrary pages of physical memory is removed, a further concern remains. Currently, all physical device drivers must operate in a

trusted domain. This is because a device that is capable of Direct Memory Access (DMA) operates on machine addresses. Therefore, it is possible to exploit the device to access areas of memory that belong to other guest operating systems, by performing DMA on those guests' address spaces.

The solution to this problem comes in the form of hardware support. An Input/Output Memory Management Unit (IOMMU) can be employed to perform mandatory translation of addresses used for DMA. Therefore, with some modifications, Xen may be used to enforce isolation by controlling what mappings may be made, in the same way as it enforces isolation by controlling page table updates. IOMMUs will be available in forthcoming chipsets from AMD and Intel.

3 Interaction with the Proof of Concept

3.1 Description of the PET Proof of Concept prototype

The Private Electronic Transaction (PET) Proof of Concept (PoC) prototype was presented to the Review Board in December 2006. The PoC prototype uses Xen at the hypervisor layer to provide isolation and other trusted computing facilities, as described in Section 2. In this section, the structure and functionality of the prototype is described.

The prototype supports two “user modes” chosen at boot time: expert and normal. The former is useful for understanding the internal details of the prototype while the latter make the prototype behave as in a real scenario.

The prototype is divided into three domains; in the normal mode:

1. dom0. The user does not interact with this domain; it uses automatic scripts to launch the two other domains, and configure the networking.
2. domU. This is an untrusted domain, which is used to run arbitrary, untrusted programs.
3. domT. This is a trusted domain, which is used to run the web browser for carrying out the Private Electronic Transaction with a bank.

The use case for the demonstrator is that, on a day-to-day basis, the user will run and install arbitrary code in domU. This code may include malevolent software, such as a keylogger or other program that monitors how the user uses the computer. When the user wishes to make a private electronic transaction, he wants to be sure that his details are safe from this untrusted code. Therefore, he switches to domT, which is a simple domain that contains only a web browser, and which is in a known-good configuration. DomT's configuration and its file system are measured (i.e. digested) and the measurement is used to extend a TPM register (PCR), which makes it possible to carry out remote attestation of the client by the server (i.e. the bank web site) through a pair of proxies that implement the attestation protocol. These proxies run in dom0 on the client side and as front-end for the web server on the other side.

3.2 Description of the development domain

In order to facilitate the development of the PET demonstrator (see above for more details), a development domain has been created. This is a full copy of a Debian Sarge Linux distribution, which includes the necessary tools and libraries in order to build Xen 3.0.4. This makes it possible to experiment with different configurations of Xen and the Linux kernel. The development domain can also be used to build the new-architecture Xen and domB components of this deliverable.

The development domain comprises a file system image (containing the Linux distribution) and a Xen domain configuration file. Due to the size of the development domain, it must be copied to the hard disk before it can be used.

3.3 Building the prototype using the development domain

The PET PoC prototype is a LiveCD, which may be used to run Xen without installing

any software onto a computer. In combination with the development domain (which does require some software to be copied to a hard disk, due to memory and space constraints), it is possible to build the domain builder prototype (as described in Subsubsection 2.2.1). Furthermore, since the development domain runs in a virtual machine, it is possible to run the modified version of Xen itself in a virtual machine. This is safer than running an experimental version of Xen directly on the physical hardware, because the effects of kernel misconfiguration can be contained within the virtual machine. Therefore, the risk of data loss during development is mitigated.

4 Conclusion

We have designed and developed a prototype of the new architecture for the Xen virtualisation layer. This architecture supports an initial reduction of the Trusted Computing Base and provides the basis for the implementation of the BMI service. It also enables the support of further research on improvement of access control mechanisms within the hypervisor. We also hope that the development domain created during this work will be used by other OpenTC partners (and external entities when becoming public) to build higher level applications. We anticipate this work to be of great use for Work packages 5 and 6 in particular.

5 List of Abbreviations

BMI	Basic Management Interface
DMA	Direct Memory Access
IOMMU	Input/Output Memory Management Unit
MAC	Mandatory Access Control
PET	Private Electronic Transaction
PCR	Platform Configuration Register
TPM	Trusted Platform Module
VM	Virtual Machine
VMM	Virtual Machine Monitor
vTPM	Virtual Trusted Platform Module
XSM	Xen Security Modules