

D04.1 Basic Management Interface

Project number	IST-027635
Project acronym	Open_TC
Project title	Open Trusted Computing
Deliverable type	Report

Deliverable reference number	IST-027635/D04.1/1.0 Final
Deliverable title	Basic Management Interface
WP contributing to the deliverable	WP4
Due date	Apr 2006 - M06
Actual submission date	Jun 12 2006

Responsible Organisation	HPLB HPLB (David Plaquin, Chris Dalton, Dirk Kuhlmann)
Authors	TUD (Bernhard Kauer, Carsten Weinhold) CUCL (Steven Hand, Steve Kell) RUB (Rainer Landfahrman)
Abstract	Definition of the Basic Management Interface requirements and first draft of an actual API.
Keywords	BMI, BMI requirements, API

Dissemination level	Public
Revision	1.0 Final

Instrument	IP	Start date of the project	1 st November 2005
Thematic Priority	IST	Duration	42 months

If you need further information, please visit our website www.opentc.net or contact the coordinator:

Technikon Forschungs-und Planungsgesellschaft mbH
Richard-Wagner-Strasse 7, 9500 Villach, AUSTRIA
Tel.+43 4242 23355 -0
Fax. +43 4242 23355 -77
Email coordination@opentc.net

The information in this document is provided “as is”, and no guarantee or warranty is given that the information is fit for any particular purpose.
The user thereof uses the information at its sole risk and liability.

Table of Contents

1	Introduction	7
2	Design Objectives.....	8
2.1	Context and scope.....	8
2.2	Interoperability.....	8
3	Related Work and Standardization Efforts.....	9
3.1	libvirt, an Open Source Library for Resource Management.....	9
3.2	CIM.....	10
3.3	Summary.....	10
4	Overview of Components and Data Model.....	11
4.1	Trusted Virtualization Layer and Boot-Up Services.....	11
4.2	Protection Domains.....	12
4.3	Controllers.....	12
4.4	Connections.....	13
5	Life Cycle and Configuration Management for XEN and L4.....	14
5.1	Xen.....	14
5.1.1	Boot of system and domain 0.....	14
5.1.2	Initial configuration of additional domains.....	14
5.1.3	CPU.....	15
5.1.4	Memory.....	15
5.1.5	Storage devices.....	15
5.1.6	Network devices.....	16
5.1.7	Pass-through of arbitrary PCI devices.....	16
5.1.8	Domain life-cycle management and dynamic configuration.....	16
5.1.9	Life-cycle management.....	16
5.1.10	Dynamic configuration.....	16
5.1.11	Configuration of Mandatory Access Control (MAC).....	17
5.2	L4.....	17
5.2.1	Boot of system.....	17
5.2.2	File providers.....	18
5.2.3	Loader and loader configuration files.....	18
5.2.4	CPU.....	19
5.2.5	Memory.....	19
5.2.6	Storage devices.....	19
5.2.7	Network devices.....	19
5.2.8	Life-cycle management.....	19
5.2.9	Mandatory access control.....	20
6	Integrity requirements and design.....	21
6.1	Boot Sequence and Chain of Trust.....	21
6.2	Design objectives.....	21
6.3	Static Root of Trusted Measurement.....	21
6.3.1	Starting Xen.....	22
6.3.2	Starting L4.....	22
6.4	Dynamic Root of Trusted Measurement.....	23
6.5	Protection Domain Integrity.....	23
7	Interface Specification.....	25
7.1	Interface: Management of Protection Domains (PDManagement).....	25
7.1.1	Method: getVersion()	26
7.1.2	Method: allocatePD()	26

7.1.3 Method: getController()	26
7.1.4 Method: getController()	26
7.1.5 Method: listPD()	27
7.1.6 Method: getDeviceManager()	27
7.2 Interface: PDController	27
7.2.1 Method: getPDID()	28
7.2.2 Method: setPriority()	28
7.2.3 Method: getPriority()	28
7.2.4 Method: getCurrentStatus()	29
7.2.5 Method: requestStatusChange()	29
7.2.6 Method: destroy()	29
7.2.7 Method: setupPD()	29
7.2.8 Method: setAllowedConnections()	29
7.2.9 Method: getAllowedConnections()	30
7.2.1 0Method: setConfigParam()	31
7.2.1 1Method: getConfigParam()	31
7.2.1 2Method: setRuntimeParam()	31
7.2.1 3Method: getRuntimeParam()	32
7.2.1 4Method: listRuntimeParams()	32
7.3 Interface: PDIImage	33
7.3.1 Method: getImageSize()	33
7.3.2 Method: getData()	33
7.4 Interface: IntegrityManagement	34
7.4.1 Method: getIntegrityLink()	34
7.4.2 Method: getProtectedCredentials()	34
7.5 Data Structures	35
7.5.1 Structure: ConnectionPolicy	35
7.5.2 Structure: ConnectionIntegrityPolicy	35
7.5.3 Structure: IntegrityReport	35
7.5.4 Structure: IntegrityPolicy	36
7.5.5 Structure: ProtectedCredentials	36
7.5.6 Structure: PDDescription	36
7.5.7 Structure: RuntimeCapability	37
7.6 Runtime Parameters	38
8 List of Abbreviations / Glossary	41

List of figures

Figure 1: Structure of TVI and BMI.....	12
---	----

List of Tables

Table 1: Xen domains standard configuration variables.....	14
Table 2: Directives in loader configuration files.....	18
Table 3: Important command-line options for L4Linux kernels.....	18

1 Introduction

The *OpenTC* architecture combines virtualization and Trusted Computing (TC) technology. Virtualization allows to create software compartments that act as *protection domains*, meaning that interactions between them are subjected to policy controls. Trusted Computing technology is employed to log and report integrity information about software components that are started on a system. In *OpenTC*, we use these mechanisms to ensure the integrity of the virtualization layer and its management components. A *platform* comprises the system hardware including its Trusted Computing Module (TPM) and those software components that are checked for integrity and logged into protected TPM registers.

The *Basic Management Interface* (BMI) described in this document reflects three principal requirements. First, the *OpenTC* architecture supports two different approaches to virtualization, and the interface provides a common view and interoperability from a management perspective. Second, protection domains can be created and reconfigured dynamically, and the BMI must support a minimal set of functions needed to configure and manage a hypervisor (or *virtualization layer*) and hosted components (*protection domains*) through their life cycle. Third, the BMI must cover functionalities allowing to link the integrity of the platform, as captured by Trusted Computing mechanisms, with the integrity of various other software components hosted by the hypervisors. These integrity functions are fundamental to reliably enforce security policies, and they support the implementation of security services for reporting platform integrity characteristics to remote parties.

With regard to a unified view on the different virtualization approaches, the current snapshot of the BMI focused on finding an appropriate common denominator. It is geared towards supporting the existing hypervisor prototypes and typically does not assume modifications to their current implementation. However, the work on the BMI specification has already revealed functional characteristics that are deemed to be useful in future hypervisor implementations. These extensions will be reflected in later versions of the BMI.

The document is structured as follows. In the next chapter, we outline the design objectives of the Basic Management Interface. Section 3 briefly introduces some existing and ongoing standardization work in the area of virtualization layer abstraction, and Section 4 discusses the data models that were considered for the specification of the BMI. Section 5 describes important design and management life cycles characteristics for the virtualization layers used in OpenTC, namely Xen and L4. Section 6 is dedicated to integrity management requirements for both virtualization layers and the link with the chain of trust.

Section 7 covers the actual specification of the BMI. The specification describes the programming interface in a traditional method, parameters, return value manner. The interfaces described in the BMI (such as Protection Domains and their controllers, and the Integrity Management) follow the model exposed in section 4.

2 Design Objectives

2.1 Context and scope

The Basic Management Interface is designed to provide a common view and interoperability from a management perspective, to support configuring as well as managing of hypervisors, and to expose basic functionalities to log and report the integrity of the platform and hosted components. Its design has also been driven by the goal of support functionalities required by various use cases developed in OpenTC. We anticipate the requirements of these use cases to evolve during the life time of the project, and so will the actual BMI specification. We will refer to the specific version discussed in this document as BMI v0.1.

The main focus of the BMI v0.1 was to abstract *Xen* and *L4* specific management functionality in a way that ensures minimal compatibility. The interface must allow to initialize a system in such a way that it can host applications developed for the use cases. The scope of BMI v0.1 is limited with regard to integrity functionality in that it only supports functions for reporting integrity metrics. As the use cases develop, extensions of the current BMI specification will finalize draft functions for integrity management that support secure policy enforcement by basic security services. It will also include additional functions to support more complex life cycles such as migration of protection domains, dynamic hardware configuration, and security policies.

2.2 Interoperability

A major goal of the BMI is to allow transparent interoperability for management software on both virtualization layers. In order to ensure implementability for both *Xen* and *L4*, this specification only defines an interface. In particular, it is as agnostic as possible with regard to specific architecture or programming language. The interface implementation for each virtualization layer will be described as part of future WP04 deliverables.

Due to their specific approaches to virtualization, it was not always possible to abstract away the differences of *Xen* and *L4*. In some instances, certain features proved to be impractical or even impossible to include in the current implementation of a particular type of hypervisor. As our aim for BMI v0.1 was not to impose modifications on the existing prototypic implementations of both layers, this has resulted in a small number of “visible” differences in the interface specification for both virtualization approaches. In these cases, we have taken care that the interface allows to query the exact features (aka capabilities) supported by each layer. This allows a consistent implementation of a management solution for the two types of hypervisors.

The exposure of such differences can be considered a useful input for hypervisor designs yet to be developed, and they may be addressed by a unified approach in future implementations.

3 Related Work and Standardization Efforts

As new and efficient virtualization and para-virtualization technologies are emerging, aspects of managing virtualization and hosted components are getting increasing attention. This has spawned several efforts around software engineering and standardization, both within the Open source community and the business world. *OpenTC* is primarily interested in a potential exploration and use of open source solutions and public standards. In the following, we outline two activities that have potential relevance for our efforts.

3.1 *libvirt*, an Open Source Library for Resource Management

Libvirt is an open source C library aimed at managing virtualized resources for Linux based systems. It is mainly geared towards managing virtualized Linux instances with focus on *Xen* as hosting environment.

Documentation and source code for this library is hosted at <http://www.libvirt.org>. According to this documentation, the design goal of the *libvirt* library is to implement "...building blocks for higher level management tools and for applications focusing on virtualization of a single node (the only exception being domain migration between node capabilities which may need to be added at the libvirt level)". In the following, we give a brief outline of the features and usability aspects of this library.

The stated objective of *libvirt* is to define an API that "allows to do efficiently and cleanly all the operations needed to manage domains on a node". The API focuses on single node management and excludes high level multi-nodes management features such as load balancing, etc. Nevertheless, the library will most probably include basic functionalities needed to support such management functions. *libvirt* aims to mask the heterogeneity of virtualization layers, and by providing a stable API, it intends to hide variations in future hypervisor implementations. These objectives are very close to those of WP4, in particular the abstraction of the virtualization layer.

In its current definition, however, *libvirt* suffers from several limitations which would make it difficult to use *libvirt* directly as a Basic Management Interface. First, *libvirt*'s architecture is "in a large part Xen specific since this is the only hypervisor supported at the moment". This implies that the minimal atomic entity that can be managed by *libvirt* is a typical *Xen* domain or, in other words, a complete virtualized operating system.

This is a serious limitation with regard to the *L4* architecture, since the atomic entity that can be managed for *L4* based virtualization layer is a "task". In *L4*, a complete virtualized operating system (such as *L4Linux*) is typically mapped to multiple interacting tasks. By using *libvirt*, a management application would lose the ability to control security policies between tasks composing a virtualized operating system, and would therefore have to rely on inflexible, hard coded policies. This limitation of *libvirt* would also impact design, implementation and management of security services developed in WP05, which can be implemented as entities that are much simpler than a complete virtualized operating system.

Second, the overview of the *libvirt* library explicitly states that it "won't try to provide all possible interfaces for interacting with the virtualization features". This limitation is problematic with regard to the integrity management functions of the BMI. Integrity management and TCG related functionalities will require specific interfaces and

concepts for both *Xen* and *L4*, and these interface would not be accessible by *libvirt*.

Finally, the *libvirt* interface relies on a complex computing environment being available for the implementation of the library. In particular, it relies on parsing capabilities for XML documents and POSIX-type functionalities for file and I/O as part of the library implementation. This requirement not only poses a difficulty for an L4 environment which has a limited computing environment at start up. It can also be considered as a potential weakness for the security of the system: the need for such a complex environment increases the size and complexity of the code that needs to be trusted on the managed system.

As a conclusion, while a *libvirt* type of interface might be well suited for managing virtualized operating systems at a WP05 level, the Basic Management Interface requires a lower-level, virtualization independent, secured and TCG aware definition.

3.2 CIM

The Common Information Model (CIM) is a standardization activity within the Distributed Management Task Force (DMTF). DMTF itself has more than 3500 active participants across 39 countries and spread across nearly 200 organizations. DMTF is the organization leading the development of management standards targeting interoperability for enterprise and Internet environment. It addresses the standardization of IT management in a platform-independent and technology independent manner, focusing on interface definitions, data models, communication protocols and architecture definition.

CIM defines an object model for all the resources present in a computing infrastructure. The CIM standard is expressed through schemas of UML definitions (available as MOF or XML formats) and specifications (available as descriptive text files). The latest schema publicly available at the time of the writing of this document is the version 2.12 published in April 2006. CIM schemas define information models ranging from physical resources such as network card and storage components to applications such as databases and web services. CIM also covers modelling of objects like systems (software based such as operating systems or hardware based such as servers), users, events, security and policies. Since CIM also addresses virtualization and virtualized resources, this standard is a natural reference for the Basic Management Interface definition effort.

Similarly to *libvirt*, CIM abstracts virtualized operating systems using the notion of compartment objects. As discussed earlier, this means that this level of abstraction is slightly too high level for the BMI definition. Furthermore, the models around security defined by CIM are based around “traditional” access control security. For instance, they do not cover Trusted Computing type of mechanisms, or aspects of system integrity. As a consequence CIM can not express security policies with access control based on the integrity of the virtualization layer and critical security services.

3.3 Summary

Libvirt and CIM both provide a partial basis for the work on the BMI, but the current limitations of those models are potential inhibitors success of WP04. We will follow these activities, may include some of their concepts where appropriate, and check their applicability for higher level management. For BMI v01, however, we decided to focus on the specific requirements of WP04 and OpenTC for the interface definition.

4 Overview of Components and Data Model

Each physical machine or *node* contains a set of CPUs, a hardware Trusted Platform Module (TPM), a certain amount of memory (RAM) and a set of physical devices such as network cards, graphic cards, storage interfaces (IDE, SCSI, SATA, etc...) and various other controllers. These components have traditionally been presents in physical machines and can be detailed e.g. in terms a CIM specification.

An object oriented design is the most common and appropriate approach to model computing system (CIM uses such an object oriented model for instance).

Consequently, the BMI also defines the software components, their functionalities and their relationship using an object oriented view. This has the benefit to be easier understandable and verifiable by humans. It should be noted, however, that an object oriented definition does not prescribe specific types of programming languages. Object oriented concepts can be implemented in any languages (including low level languages like assembly).

This following sections describe the principal objects manipulated in the BMI and their interactions. The detailed description of the object interfaces can be found in the section 7 of this document.

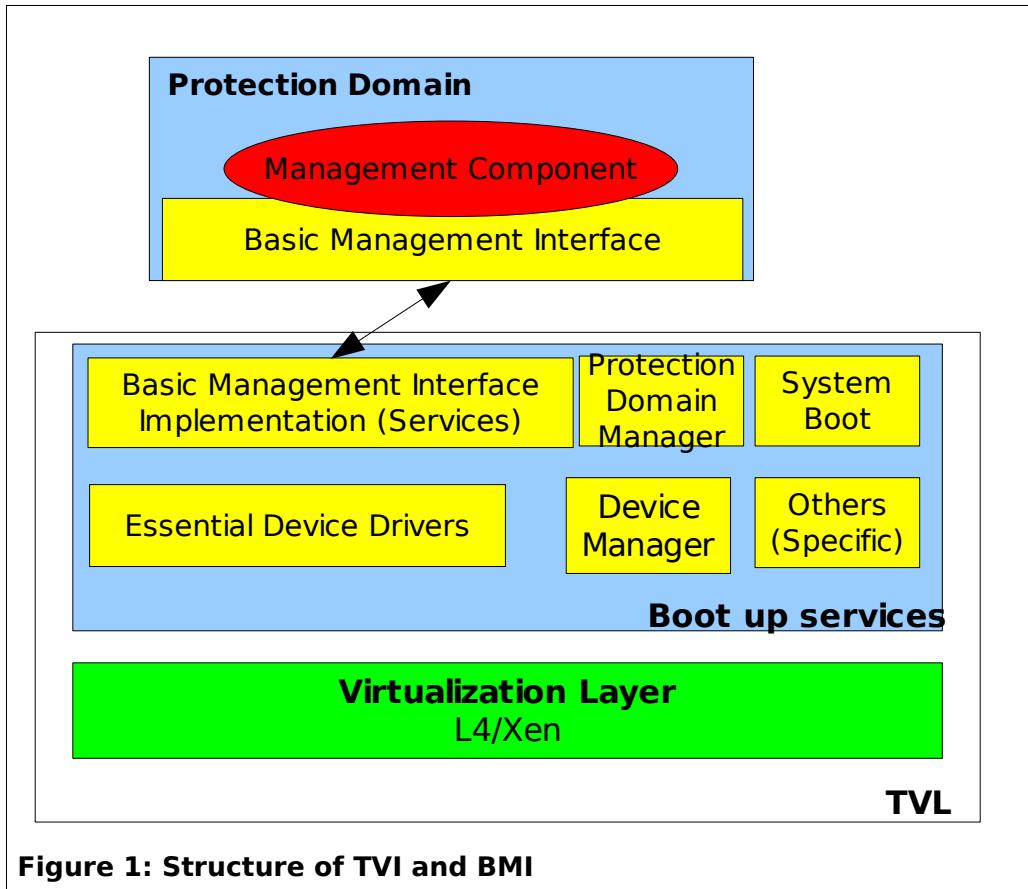
4.1 Trusted Virtualization Layer and Boot-Up Services

The implementation of the BMI relies on the existence of a Trusted Virtualization Layer (TVL). The TVL is implemented by *Xen* or *L4* as a virtualization layer, and by a set of boot-up services. These services are required to bring the physical machine to an attestable state where it can serve the requests from the management software through the BMI. They are responsible for initializing the system's hardware during system boot, for providing the TVL with essential drivers to access hardware other boot up services may rely on (such as TPM driver for instance), and for offering additional services specific to the virtualization layer.

The boot-up services also provide functionalities required for the implementation of the BMI, but they are not reflected in the BMI definition. In fact, in order to abstract the way a particular virtualization layer is booted up, the TVL is treated as a monolithic software rather than a set of manageable components. From the perspective of the BMI, the TVL is a piece of a software layer that needs to be trusted in order to allow extending this trust to the whole system.

For this reason, the TVL is also sometimes referred to as the Trusted Computing Base (TCB). However, we will avoid the expression "TCB" as the actual definition of a TCB tends to vary with regard to the application. For instance, in the case of the BMI, the TVL (together with the components of the chain of trust) is the TCB for the implementation of the BMI. On the other hand, when considering a complex scenario where a user needs to trust the whole system he is using, the TCB would comprise more services, such as the graphical user interface. A more detailed discussion on the chain of trust and integrity of the TVL is part of section 6.

While the BMI do not expose the modularity of the TVL, it defines higher level dynamic and modular "protection domains", their management interface and their interactions as seen from the point of view of the TVL. This way, management components only have to deal with the configuration and management of the dynamic part of the system once booted, and will leave the specifics of the initialization up to the TVL.



4.2 Protection Domains

Protection Domains are software components with one or more Virtual CPUs (mapped to real CPUs) and some virtual memory (mapped to the physical memory of the machine) allocated to them. They are implemented as a Domain (in Xen) or a Task (in L4). The entity which creates a new Protection Domain (such as a Management Agent) must provide a binary image describing the executable part of the Protection Domain. The format of the binary image of a Protection Domain is specific to a particular virtualization layer, but the entity using the BMI does not need to be aware of the differences. It can simply consider the binary image as an opaque piece of data associated with a Protection Domain.

Once resources have been allocated for a new Protection Domain, the BMI creates a corresponding execution environment. The TVL then schedules the Protection Domain, with a priority assigned to it through the BMI. Each Protection Domain can provide specific services to other Protection Domains. For instance, each security service implemented by WP05 could be a Protection Domain providing an interface to virtualized operating systems or other security services.

4.3 Controllers

The BMI provides a controller interface for each Protection Domain being created. The Controller allows to manage the state and configuration of a Protection Domain by

dedicated components both at creation and run time over its life cycle. The interface exposed by the controllers is therefore key in managing the Protection Domains hosted by the TVL. A protection domain can access its own controller as well as controller of Protection Domains it has created and/or has the appropriate privileges to do so.

Controllers enable the setup of Protection Domains by abstracting the source of their binary images, and by delegating the responsibility of gathering and providing the actual bits of data to the application client of the BMI. While a binary image could exists either as a file on a hard disk, as a file on a remote server, or as a stream being generated on the fly by a building component, the Controller will simply receive the binary image data and build the corresponding domain transparently.

Furthermore, Controllers include functionality allowing to set and change parameters associated to the Protection Domain. These parameters will typically be passed to the corresponding Protection Domain. Depending on the requirements of the underlying virtualization architecture, there may be cases where parameters are interpreted directly by the BMI.

Finally, controllers provide access to the Integrity Management interface. This interface creates the link between the integrity of the TVL and the integrity of the controlled Protection Domain (see section 6 for a more detailed discussion on integrity).

4.4 Connections

Protection Domains need to interact with each others to provide the services they implement. For instance, a Protection Domain running a virtualized operating system has to interact with Protection Domains providing device drivers for network, storage, potentially keyboard, etc.

These interactions take place over communication channels that are referred to as *Connections* in the BMI definition. The implementation of the communication channels is specific to the virtualization architecture, and the actual semantic of the messages exchanged over those channels depends on the type of the implemented services. For instance, Xen relies on a protocol based on shared memory and events to implement communication channels, while L4 uses an message based Inter Process Communication (IPC) approach. This is the main reason why the BMI currently does not specify how to establish connections between Protection Domains, or how to pass messages through Connections. As the semantic and format of messages vary for each service, they will be defined by the teams working on the virtualization layers.

Although the BMI does not specify an interface for establishing and using communication channels, it provides functionalities to control access to connections. The access control policies will be enforced by the TVL when Protection Domains attempt to create or use a communication channel. A management component can therefore control the topology of the allowed communications between the Protection Domains under its control.

As the trustworthiness of a system may not only depend on its own integrity, but also on that of other systems it can interact with, the Management Interface of the BMI supports basic functionalities for combining connection policies with system integrity. This allows to implement applications that can remotely attest to the integrity of the TVL it runs on, as well as all the other Protection Domains it depends on for its implementation.

5 Life Cycle and Configuration Management for XEN and L4

This section describes the existing processes under *Xen* and *L4* involved in creating and configuring a Protection Domain (referred to as a *domain* for *Xen* and a *task tree* for *L4*). These descriptions are used to validate the functionalities of the BMI.

5.1 Xen

The following paragraphs outline the life cycle of a domain in *Xen*, including how it receives configuration information about devices it may access and how access control policy fits in.

- boot of host machine, *Xen, Domain 0*
- pre-configuration of domains
- life-cycle management of domains (incl. stopping, rebooting)
- dynamic configuration of domains
- configuration of mandatory access control (MAC) policies and mechanisms

5.1.1 Boot of system and domain 0

The first domain, *Domain 0*, is created automatically when the system boots up. This domain has special management privileges: it builds other domains and manages their virtual devices. It also performs administrative tasks such as suspending, resuming and migrating other virtual machines. Within *domain 0*, a daemon process called *xend* is responsible for managing virtual machines and providing access to their consoles. Commands are issued to *xend* over an HTTP interface, via a command-line tool called “*xm*”.

5.1.2 Initial configuration of additional domains

Before an additional domain can be started, a corresponding configuration must be created. *Xen* configuration files contain the following standard variables.

Table 1: Xen domains standard configuration variables

Variable	Description
kernel	Path to the kernel image.
ramdisk	Path to a ramdisk image (optional).
memory	Memory size in megabytes.
vcpus	The number of virtual CPUs.
root	Specify the root device parameter on the kernel command line.
nfs_root	Path of the root filesystem on the NFS server (if any).
nfs_server	IP address for the NFS server (if any).

Variable	Description
Disk	List of block devices to export to the domain, e.g. <code>Disk = ['phy:hdal, sdal, r']</code>
hostname	Set the hostname for the virtual machine.
dhcp	Set to "dhcp" if you want to use DHCP to configure networking.
netmask	Manually configured IP netmask.
gateway	Manually configured IP gateway.
vif	Network interface configuration. This may simply contain an empty string for each desired interface, or may override various settings, e.g. <code>vif = ['mac=00:16:3E:00:00:11, bridge=xen-br0', 'bridge=xen-br1']</code>
console	Port to export the domain console on (default 9600 + domain ID).
extra	Extra string to append to the kernel command line (if any)

Additional fields (e.g. to configure virtual TPM functionality) are documented in example configuration files included with the Xen documentation in the /etc/xen directory of Xen's distribution. When created, a domain is assigned an ID by the hypervisor which will then be used to reference it in later system calls.

5.1.3 CPU

Xen allows a domain's virtual CPU(s) to be associated with one or more host CPUs. This can be used to allocate real resources among one or more guests, or to make optimal use of processor resources when utilizing dual-core, hyperthreading, or other advanced CPU technologies.

5.1.4 Memory

Memory may be partitioned statically; alternatively, domains allow dynamism by implementing a ``balloon'' driver. This is a driver which occupies a variable amount of a domain's memory, and gives that memory back to Xen. The result is that the guest operating system is under the illusion that it has a fixed amount of physical memory, but in fact, the area being occupied by the balloon driver is pooled for use by other domains. The size of the balloon may be dynamically altered, but can never cause the domain to occupy more memory than its configured maximum; there is also a corresponding minimum when the balloon is full.

5.1.5 Storage devices

Storage can be made available to virtual machines in a number of different ways. The most straightforward method is to export a physical block device (a hard drive or partition) from *Domain 0* directly to the guest domain as a virtual block device (VBD). Storage may also be exported from a filesystem image or a partitioned filesystem

image as a file-backed VBD. Finally, standard network storage protocols such as NBD, iSCSI, NFS, etc., can be used to provide storage to virtual machines.

5.1.6 Network devices

Each domain network interface is connected to a virtual network interface in *Domain 0* by a point to point link (effectively a virtual crossover cable). Traffic on these virtual interfaces is handled in *domain 0* using standard Linux mechanisms for bridging, routing, rate limiting, etc. *xend* calls on two shell scripts to perform initial configuration of the network and configuration of new virtual interfaces. By default, these scripts configure a single bridge for all the virtual interfaces. Arbitrary routing / bridging configurations can be configured by customizing the scripts.

5.1.7 Pass-through of arbitrary PCI devices

Individual PCI devices can be assigned to a given domain (a PCI driver domain) to allow that domain direct access to the PCI hardware. The PCI devices you wish to assign to unprivileged domains must be hidden from your backend domain (usually domain 0) so that it does not load a driver for them. This is done using the *pciback.hide* kernel parameter which is specified on the kernel command-line and is configurable through *GRUB*. An example kernel command-line which hides two PCI devices might be:

```
root=/dev/sda4 ro console=tty0 pciback.hide=(02:01.f)(0000:04:1d.0)
```

PCI devices can also be bound and unbound to the PCI backend after boot through the manual binding/unbinding facilities provided by the Linux kernel in *sysfs*.

5.1.8 Domain life-cycle management and dynamic configuration

Life-cycle management functions are exposed through *xend*'s HTTP server interface, and may be performed from the command-line using the *xm* tool. Various dynamic configuration options are also available.

5.1.9 Life-cycle management

The interface exposed by *xend* provides operations that can be used to

- create, pause, shutdown, reboot and destroy domains
- assign friendly name (basically a string) to created domains
- save and restore domains (to and from files in *domain 0*)
- migrate a domain to a different physical host, in either a live or non-live manner
- send system request like signals to domains (Linux domains only)

5.1.10 Dynamic configuration

In addition to life-cycle management, *xend* exposes the following configuration operations:

- change the maximum and current memory allocations of a domain
- enable virtual CPUs (VCpus) in a domain (up to the configured maximum)
- pin a VCPU to a physical CPU

- attach, detach or list virtual block devices
- attach, detach or list virtual network devices
- create, destroy or list the host's virtual networks

As well as the operations exposed by *xend*, dynamic binding and unbinding of PCI devices can be performed using the Linux *sysfs* interface in *domain 0*.

Note that *xend* also supports the following monitoring and diagnostic operations:

- print the *Xen* message buffer or the *xend* log
- list current domains, and convert between friendly "names" and IDs
- display a real-time domain resource monitor (*xentop*)

5.1.11 Configuration of Mandatory Access Control (MAC)

MAC policies are binary blobs compiled from XML by a tool. Typical policies are Type Enforcement and Chinese Wall Policies. They are normally loaded by the bootloader as a separate module. Additionally, each domain and resource is assigned a security label, which will be used to make an access control decision by the policy module. Some additional *xm* commands are provided for live manipulation of these policies. These allow:

- dynamic loading of policy (currently very limited: only works when an unrestricted default policy is in place, or if the new policy has the same name as the existing one)
- listing the labels defined by a particular policy
- displaying the label of each domain (in addition to other fields)
- applying a label to a domain.

The configuration file syntax has also been extended to include an 'access_control', option which may specify a policy and label for a domain.

Finally, there are separate tools provided for the following operations:

- get an access control decision for a pair <domain, label>
- get the label for a given domain.

5.2 L4

In the *L4* case, domain-like virtual machines are provided by instances of L4Linux 2.6. The L4Linux kernel itself is an L4Env application that runs in an *L4* task. All Linux applications running within an L4Linux instance are implemented as *L4* tasks as well, forming a *task tree* with the kernel at the root. Individual L4Linux instances are started by an L4Env service called *loader*, which is also given the configuration for the newly started virtual machine.

5.2.1 Boot of system

The *L4* microkernel and the basic L4Env services are loaded and bootstrapped by the boot loader. Binaries and initial configuration files can be read from machine-local storage (e.g., a CD-ROM or hard disk accessible by the boot loader). Another

alternative is to retrieve all binaries and configuration files required during boot up from a TFTP server on the network.

5.2.2 File providers

Once the microkernel and the basic services are running, the loader service can be used to start new L4Linux instances. L4Linux kernel binaries, initial RAM disks, and loader configuration files are retrieved via the file provider interface as defined by the *generic_fprov* package. The following implementations of this interface are available:

- **bmodfs**: This is a service that offers files provided by the boot loader *GRUB*. Any files to be offered have to be passed to *bmodfs* as *GRUB* modules. These modules must be specified in the *GRUB* configuration file *menu.lst* and will then be loaded into system memory during boot up of the host system.
- **tftp**: This service provides files stored on a TFTP server running on another machine on the network. It has built-in network drivers and can also use ORe, a network multiplexer that allows it to share a single network interface card with, for example, multiple instances of L4Linux.

5.2.3 Loader and loader configuration files

L4Linux instances are configured at start-up time via information provided in *loader* configuration files. In order to start an L4Linux instance, the *loader* interprets the specified configuration file. Table 2 lists configuration directives relevant to starting new L4Linux instances.

Table 2: Directives in loader configuration files

Directive	Description
task	Location of the kernel image and command-line options.
module	Additional modules such as a ramdisk image (optional).
file_provider	The file provider from which the kernel and ramdisk images are to be loaded.
priority	L4 scheduling priority.
ioport	I/O ports to which access is allowed.
allow_cli	Allow to disable interrupts.

Additional configuration options are passed to the L4Linux kernel on its command line as specified using the *task* directive in the configuration file. Table 3 gives an overview of command-line options commonly used for L4Linux kernels.

Table 3: Important command-line options for L4Linux kernels

Variable	Description
mem	Size of main memory available to the L4Linux instance.

Variable	Description
load_ramdisk	Load a ramdisk.
ramdisk_size	Maximum size of ramdisk images.
l4env_rd	Location of the ramdisk image.
video	Used to specify options for the L4 framebuffer driver.

5.2.4 CPU

The *L4/Fiasco* microkernel and L4Linux instances running on top of it can use only one CPU. This CPU is shared among multiple L4Linux instances.

5.2.5 Memory

Memory is allocated statically to L4Linux instances using *loader* configuration files. The L4Linux kernel makes parts of this memory available to the *L4* tasks that represent its Linux processes.

5.2.6 Storage devices

Persistent storage can be made available for any L4Linux instance. One of the instances can access the hard disk controller and the attached disks directly. Others can use standard network storage protocols such as NFS and iSCSI.

5.2.7 Network devices

L4Linux instances can access the network through a stub kernel driver for the *ORe* network multiplexer. *ORe* allows multiple clients (i.e., L4Linux instances) to share one physical network interface card. It also enables its clients to communicate locally using network protocols. Multiple instances of *ORe* can also be used to create isolated vLAN segments. A dedicated instance of L4Linux connected to these vLANs can then act as a router and firewall.

5.2.8 Life-cycle management

To control the life cycle of L4Linux instances, we provide a light-weight *L4* application called *run*, which can be accessed via the *L4* graphical console *l4con*. The *run* application provides interfaces to:

- start applications using the *loader* service
- kill applications
- show run time information
- dump memory
- enter the kernel debugger of *L4/Fiasco*
- reboot the host machine

Additionally, two *L4*/Linux hybrid programs, *run-l4* and *kill-l4*, can be used from within an L4Linux instance to start or stop other instances of L4Linux. Just like *run*, *run-l4* also uses the loader service and loader configuration files in order to start a new instance of L4Linux.

Functionality to support suspending and resuming of L4Linux instances is currently being developed.

5.2.9 Mandatory access control

Currently, the L4Fiasco microkernel does not enforce restrictions regarding communication (IPC) among *L4* tasks. Server tasks and applications have to implement access control and the according policies by themselves. As a short-term solution for IPC control we will provide an implementation of a reference monitor that can restrict IPC among communication partners. It will be capable of allowing or disallowing communication between specific *L4* tasks, and can therefore be used to enforce information flow policies. The specific data formats and interfaces have yet to be determined.

6 Integrity requirements and design

Trusted Computed technology, as specified by the Trusted Computing Group (TCG), describes how to measure, report and use a system integrity to either attest it to a remote party or to locally enforce policies. However, TCG addresses the integrity of the system from the point of view of its boot sequence and specifies how to establish a chain of trust of the components involved during the boot process. The integrity of dynamically allocated and created components, such as the Protection Domains of the BMI, is not part of the scope of the TCG specification. The goal of the Integrity related mechanisms of the TVL, exposed through the BMI, is to provide a link between the integrity of the system and its TVL (established by the chain of trust) and the various Protection Domains running on the TVL.

6.1 Boot Sequence and Chain of Trust

This section gives an overview of the boot sequence and measurement sequences executed to implement the Chain of Trust for each of the 2 virtualization layers. It mainly focuses on examining what is measured and stored into the Platform Configuration Registers (PCRs) of the TPM during the boot of Xen and L4 respectively.

6.2 Design objectives

Integrity measurements, as defined by TCG, cryptographically depend on the actual binary data of the executed software components during the boot, and also depend on the exact order in which they are executed. As a result, in order to limit the heterogeneity of integrity measurements obtained after boot up on various systems, and in order to simplify software update operations, the design of the integrity mechanisms of the TVL aims at reducing the number of measured software components as much as possible. Implementations of the TVL and its chain of trust will therefore try to minimize the number of boot up services part of the measured sequence, while ensuring that all the critical ones are part of the measured sequence in order not to introduce any security weakness in the system.

The integrity measurement of the boot sequence will start from the Root of Trust for Measurement (RTM) and will carry on until the TVL and all its security sensitive boot up services are measured. The integrity of the TVL (and of the underlying system) will then serve as the basis for the reporting of the integrity of the Protection Domains. In order to maintain the overall security of the system and the isolation requirements between Protection Domains, the TVL will ensure that integrity of Protection Domains is independent from each other. This is also true for the enforcement of security policies based on the integrity of individual compartments.

6.3 Static Root of Trusted Measurement

This section describes the sequence of events that happens during a boot of a system when using the static root of trust mode. In this mode, the Core RTM (CRTM) is composed of the CPU and the Bios Boot Block (BBB) which is the first piece of code executed after a power on (or a reboot) of the system. The system is therefore in a well known state (reset) and the first executed code is also known. The sequence of events is the following:

1. The BBB measures the BIOS and extends the appropriate PCR with this value

2. The BIOS measures all its Option ROMs and extends the appropriate PCR with this value
3. The BIOS measures the Master Boot Record (MBR) of the booted media and extends the appropriate PCR with this value
4. The MBR (which is part of Stage1 of *Trusted GRUB*) measures the first sector of the stage2 file and extends the appropriate PCR. It then transfers control to the first sector of stage2.
5. The first sector of Stage2 loads and measures the rest of the stage2 file. It then extends the appropriate PCR with the measured value.
6. Stage2 loads the configuration file and waits for either a default choice to happen or for the user to execute a command.
7. Once the user has select one configuration and validated its choice, or once a default configuration has been chosen due to a timeout, *Trusted GRUB* measures the “string” composed of the commands and their parameters, and extends the appropriate PCR with this value.
8. *Trusted GRUB* loads every module into memory and measures them individually. It then extends the appropriate PCR with those values.
9. Execution is transferred to the first loaded module.

6.3.1 Starting Xen

In the case of Xen, the modules loaded by Trusted GRUB are:

1. Binary image of Xen executable
2. Kernel image of Domain 0
3. An optional initial ramdisk used by Domain 0.

Trusted GRUB will transfer execution to Xen executable image at stage 9 in the boot sequence. Once Xen has finished to initialize it will then start to schedule and execute the code contained in Domain 0. This privileged domain will then be able to carry on the boot sequence of the system. Depending on the architecture of this Domain, some other measurement might be carried out by Domain 0 and reported to the TPM. Indeed, Domain 0 is part of the implementation of the TVL with regards of the BMI. So, if Domain 0 includes modules that are loaded from a source external to the initial ramdisk, it will have to ensure the integrity of these modules is measured and included in the integrity of the TVL.

6.3.2 Starting L4

In the case of L4, the modules loaded by Trusted GRUB are:

1. *L4 MicroKernel image*
2. *Startup modules such as Root task, task manager, task loader, memory manager, IO Manager, TPM Driver, and VTPM in a L4Env configuration*

Similarly to the case of Xen, Trusted GRUB will transfer execution to the root task at stage 9 in the boot sequence. The root task will then be used to start some other tasks. The integrity requirement is similar to the previous case if the root task need to

rely on some other “external” (i.e. Not measured by Trusted GRUB) tasks to implement the BMI.

6.4 Dynamic Root of Trusted Measurement

This section describes the sequence of events that happens when a system is booted using the Dynamic root of trust mode. In this mode, the Core RTM (CRTM) is composed of the CPU, and the chipset of the system. The dynamic mechanism allows the first piece of code to be measured by the CPU directly and reported to the TPM. The CPU and the chipset of the system is placed in a well known state, and the chain of trust then starts with the first software module. This method requires an extra software module (Open Source Loader, OSLO) as part of the chain of trust which is used to start the dynamic measurement process and boots the virtualization layer.

1. *GRUB*, or another application if running in a traditional operating system, sets up the memory to receive the virtualization layer code and various other modules. This is similar to step 8 of the static RTM case, except that no measurement are sent to the TPM.
2. *OSLO* is then started by either *GRUB* or the other application.
3. *OSLO* then performs a CPU specific instruction which starts the dynamic chain of trust process. On AMD processors, this instruction is called “skinit”
4. The CPU measures the initial boot code present in *OSLO*.
5. *OSLO* then measures all the modules based on the Multi Boot Information (MBI) data passed by *GRUB*, and extend the appropriate PCR with this measurement. This stage is similar to the measurement done by Trusted GRUB in step 8 of the Static RTM case.
6. *OSLO* then starts the first module the same way as Trusted GRUB in step 9 of the Static RTM.

After this step, the remaining of the boot sequence for Xen and L4 is the exactly the same as the one described in the case of the Static RTM.

6.5 Protection Domain Integrity

Once the TVL and all its boot up services have been started, the TVL can serve requests from the BMI to create new Protection Domains. The TVL will provide a link between its own integrity and the integrity of the new Protection Domains by performing a measurement of the Protection Domain binary image. Because the format of that image is specific to each virtualization layer, the TVL is responsible for parsing the image appropriately in order to determine what portions of it need to be part of the measurements. This measurement will happen before any part of the Protection Domain gets executed.

The TVL will then keep this measurement in a memory area inaccessible to the measured Protection Domain or any other Protection Domain which was not involved to create the measured Protection Domain. This requirement comes from the need for isolating the integrity of the Protection Domains as explained in the Design Objectives section.

Access to the measurement will only be allowed through the Integrity Management interface of the BMI. Through this interface, a Protection Domain can obtain a digital

signature form the TVL of the measurement of its integrity together with the content of the PCRs of the TPM. Such a signature will prove to a remote party the integrity of the whole chain of trust and of the Protection Domain itself. The Protection Domain can also use the Integrity Management to recover credentials protected by an Integrity Policy. In this case, the TVL will enforce that the Protection Domain can only have access to the credentials (or any kind of secret) if the integrity measurement stated in the Integrity policy matches the current Integrity measurement of the Protection Domain.

By designing security services that use this mechanisms, it should be possible to implement a system that enforces security policies of a higher level of abstraction.

7 Interface Specification

This section describes interfaces provided by the virtualization layer and implemented by an abstraction library. Typical software clients making use of this interface are management components and high level security services as developed by WP5. We anticipate that there will be library implementations that are specific for Xen and L4, respectively. The library interface will provide interoperability for the management software running on a particular virtualization layer. The functions exposed by this interface can be implemented as integral parts of the virtualization layer or, alternatively, as part of a minimal set of startup services running on top of this layer. Users of this interface should not make assumptions on whether it is linked to functions provided by the virtualization layer or by startup services. The interface specification was originally written in XML format. This allows flexibility and extensibility regarding properties of methods and data structures. Furthermore, this approach is implementation language agnostic. At this stage, we have implemented translation into HTML to facilitate the review process. In future, we will also provide translations into specification and programming languages such as IDL, C/C++ and Java. The current specification of the WP4 Management Interface should not be considered as stable. It is the first snapshot of an iterative process. It will evolve during the lifetime of the OpenTC project as new functionalities requirement are identified.

7.1 Interface: Management of Protection Domains (PDManagement)

This interface regroups management functions as exposed by the very first layer of software to manage the hypervisor. It is assumed that the entity using this interface is privileged to use it. Therefore, the authentication and access control mechanisms are not part of this interface definition. This interface can be provided in the form of a library which implementation is specific to the underlying virtualization layer. At its lowest level, the interface allows to manage allocatable resources called "Protection Domains". A Protection Domain is an executable component that receives a portion of memory and of CPU cycles and is scheduled by the hypervisor. It is an abstract component in that it can, for example, refer to L4 tasks, virtual devices, compartments, and services. It typically provides a service to other Protection Domains through a communication channel. We assumed that all platform management requests will be initiated through this interface only, allowing to make the implementation of the interface stateful if needed.

The PDManagement interface includes the following methods:

- PDManagement.getVersion()
- PDManagement.allocatePD()
- PDManagement.getController()
- PDManagement.getController()
- PDManagement.listPD()
- PDManagement.getDeviceManager()

7.1.1 Method: getVersion()

- **Function:** Return the version of the API implemented by this interface.
- **Parameters:** NONE
- **Return value:**

Name	Type	Description
	int	The API version implemented by the library/API. The 16 Most Significant Bits contain the major version and the 16 Less Significant Bits contain the minor version. Latest version is 0x00000001.

7.1.2 Method: allocatePD()

- **Function:** Allocate resources to create a new Protection Domain. After a call to this method, the newly created Protection Domain is not running but should be ready to run.
- **Parameters:**

Name	Type	Description
desc	[in] PDDescription	The resource requirement for the new allocated Protection Domain.

- **Return value:**

Name	Type	Description
	PDController	Reference to the controller for the allocated Protection Domain, or a NULL reference if the operation failed.

7.1.3 Method: getController()

- **Function:** Get a controller for the Protection Domain with the local identifier **id**.
- **Parameters:**

Name	Type	Description
id	[in] int	Local Identifier of the Protection Domain to retrieve a controller for.

- **Return value:**

Name	Type	Description
	PDController	Reference to the controller of the Protection Domain with Local Identifier id , or NULL if there is no Protection Domain with this id .

7.1.4 Method: getController()

- **Function:** Request the controller of the current Protection Domain, that is, the Protection Domain calling this method.

- **Parameters:** NONE

- **Return value:**

Name	Type	Description
	<i>PDController</i>	Reference to the controller of the calling Protection Domain. This method should never return NULL.

7.1.5 Method: listPD()

- **Function:** Lists the ID of all Protection Domain currently allocated by the hypervisor.

- **Parameters:** NONE

- **Return value:**

Name	Type	Description
	<i>int[]</i>	List of Local ID of currently allocated Protection Domain

7.1.6 Method: getDeviceManager()

- **Function:** Get an interface to configure the devices of the system.

- **Parameters:** NONE

- **Return value:**

Name	Type	Description
	<i>DeviceManager</i>	A reference to the device manager of the system.

7.2 Interface: PDController

This interface allow to set up and configure a Protection Domain.

The PDController interface includes the following methods:

- PDController.getPDID()
- PDController.setPriority()
- PDController.getPriority()
- PDController.getCurrentStatus()
- PDController.requestStatusChange()
- PDController.destroy()
- PDController.setupPD()
- PDController.setAllowedConnections()
- PDController.getAllowedConnections()

- PDController.setConfigParam()
- PDController.getConfigParam()
- PDController.getRuntimeParam()
- PDController.getRuntimeParam()
- PDController.listRuntimeParams()

7.2.1 Method: getPDID()

- **Function:** Get the Local ID of the Protection Domain controlled by this interface. The ID is local to the physical machine and specific to the hypervisor's implementation.
- **Parameters:** NONE
- **Return value:**

Name	Type	Description
	int	The ID of the Protection Domain controlled by this interface. Note that there is no guarantee on the current status of the Protection Domain. Also the Protection Domain controlled by this interface might not exist anymore. The only way to know the status of a Protection Domain is to call the see: getStatus() method.

7.2.2 Method: setPriority()

- **Function:** Changes the priority of the Protection Domain. The underlying implementation might not support changing the priority of a Protection Domain at run-time.
- **Parameters:**

Name	Type	Description
priority	[in] int	Relative priority in the range of -10 to +10. Priority 0 is the default priority for a Protection Domain.

- **Return value:** NONE

7.2.3 Method: getPriority()

- **Function:** Get the current priority of the Protection Domain controlled by this interface.
- **Parameters:** NONE
- **Return value:**

Name	Type	Description
	int	The current priority of the Protection Domain

7.2.4 Method: getCurrentStatus()

- **Function:** Gets the current status of the Protection Domain.
- **Parameters:** NONE
- **Return value:**

Name	Type	Description
	PDStatus	The current status of the Protection Domain. see: PDStatus type for more information.

7.2.5 Method: requestStatusChange()

- **Function:** Request the Protection Domain to switch to a new status
- **Parameters:**

Name	Type	Description
status	[in] PDStatus	The new status the Protection Domain should switch to. Supported Status are implementation dependent, so a client of this interface should not assume the operation was sucessful.

- **Return value:** NONE

7.2.6 Method: destroy()

- **Function:** Destroys the Protection Domain controlled by this interface and free all resources the Protection Domain was using.
- **Parameters:** NONE
- **Return value:** NONE

7.2.7 Method: setupPD()

- **Function:** Instruct the virtualization layer to build the Protection Domain content. On executing this command, the virtualization layer will use the provided PDIImage **desc** interface to build the memory image of the Protection Domain.
- **Parameters:**

Name	Type	Description
desc	[in] PDIImage	Client provided interface for fetching the Protection Domain data.

- **Return value:** NONE

7.2.8 Method: setAllowedConnections()

- **Function:** Allow a connection to take place between the Protection Domain represented by the PDController instance and another Protection Domain on the same platform referred to by the local id **tid**. . The connection will be used by the controlled Protection Domain to require services from the Protection Domain identified by **tid**. By default connections are not allowed unless specifically

allowed by this method. Protection Domains will then use the usual implementation specific interface provided by the Hypervisor to open connections to other Protection Domains. This operation will succeed only if the connections were specifically allowed using this method. Note: Previous allowed connection to the Protection Domain

- **Parameters:**

Name	Type	Description
<i>tid</i>	[in] int	The Local ID of the Protection Domain to which the connections policy need to be set.
<i>allowed_connections</i>	[in] <i>ConnectionPolicy[]</i>	A list of communication type allowed for this connection. Note: The underlying implementation might only support a subset of the types. Clients of this interface should check the supported connection types prior to this call.
<i>disable_existing</i>	[in] boolean	Specifies if existing connections should be disabled (if they are not part of the allowed_connections parameter). If true, existing connections that are not explicitly allowed by the allowed_connections parameter will be closed. If false, the connections will remain open until the Protection Domain closes them itself using the standard hypervisor API.

- **Return value:** NONE

7.2.9 Method: **getAllowedConnections()**

- **Function:** Get the list of connections allowed from the Protection Domain controlled by this interface to another specific destination Protection Domain.

- **Parameters:**

Name	Type	Description
<i>tid</i>	[in] int	The Local ID of the Protection Domain to which the connections policy relative to the Protection Domain controlled by this interface need to be retrieved.

- **Return value:**

Name	Type	Description
	<i>ConnectionPolicy[]</i>	Array of allowed connection types from this Protection Domain to the Protection Domain tid .

7.2.10 Method: setConfigParam()

- **Function:** This method sets configuration parameter for the Protection Domain being created. The created Protection Domain will use the corresponding getConfigParam() function to retrieve its configuration at start up. Some of the configuration parameters may then be exposed as Runtime Parameters if they can be altered later on.

- **Parameters:**

Name	Type	Description
name	[in] string	The name of the parameter to set. see: RuntimeCapability for more details about the naming conventions.
value	[in] string	The value of the parameter to set. Note: It is the responsibility of the caller to translate the value of non string parameter into a string representation

- **Return value:** NONE

7.2.11 Method: getConfigParam()

- **Function:** Allow a Protection Domain to retrieve a parameter passed to it by the instance that created it.

- **Parameters:**

Name	Type	Description
name	[in] string	The name of the parameter to retrieve. see: RuntimeCapability for more details about the naming conventions.

- **Return value:**

Name	Type	Description
	string	The value of the requested parameter. Note: It is the responsibility of the caller to translate the string representation of the parameter into the correct data value of the expected type.

7.2.12 Method: setRuntimeParam()

- **Function:** Modifies a Parameter of the currently running Protection Domain.

- **Parameters:**

Name	Type	Description
name	[in] string	The name of the parameter to set for the controlled Protection Domain. The caller should not assume a parameter exists. It should first check for the existence of a parameter using the listRuntimeParams() function. see:

Name	Type	Description
value	[in] string	Description RuntimeCapability for more details about the naming conventions. The value of the parameter to set. Note: It is the responsibility of the caller to translate the value of non string parameter into a string representation

- **Return value:** NONE

7.2.13 Method: getRuntimeParam()

- **Function:** Retrieve the current value of a Parameter for a Protection Domain. Protection Domain may chose to expose their status and capabilities through the use of these dynamic parameters.
- **Parameters:**

Name	Type	Description
name	[in] string	Description The name of the parameter to retrieve for the controlled Protection Domain. The caller should not assume a parameter exists. It should first check for the existence of a parameter using the listRuntimeParams() function. see: RuntimeCapability for more details about the naming conventions.

- **Return value:**

Name	Type	Description
	string	Description The current (dynamic) value of the parameter name .

7.2.14 Method: listRuntimeParams()

Name	Type	Description
root	[in] string	Description The parent parameter to list the descendant parameters from. For instance, using "/supported-capabilities" will return a list of parameters named "/supported-capabilities/****". If "/" is specified, this will list the first level of parameters. Note: Specifying the empty string "" is equivalent to specify "/". Note: Any trailing "/" is ignored..

- **Return value:**

Name	Type	Description
	RuntimeCapability[]	The list of parameters child of root and their supported access mode. see: RuntimeCapability for more details.

7.3 Interface: PDImage

Abstract the provisioning of Protection Domain's data image. This allows the client of this interface to abstract the location of the image (such as file, URL, etc...). The methods of this interface will be called by the implementation of the BMI library when setting up a Protection Domain, see: **setupPD()** method. The client of this library will therefore need to provide its own implementation of this interface to allow the BMI to retrieve the binary image of the Protection Domain the client wishes to create.

The PDImage interface includes the following methods:

- PDImage.getImageSize()
- PDImage.getData()

7.3.1 Method: getImageSize()

- **Function:** Get the size of the binary image of the Protection Domain.
- **Parameters:** NONE
- **Return value:**

Name	Type	Description
	int	Size of the Protection Domain binary image in bytes.

7.3.2 Method: getData()

- **Function:** Read a set of the data image describing the Protection Domain.
- **Parameters:**

Name	Type	Description
offset	[in] int	The offset within the binary image. Note: must be within the range of 0 and the actual image size,
length	[in] int	Number of bytes to retrieve.

- **Return value:**

Name	Type	Description
	byte[]	Requested bytes of the binary images. Note: The number of bytes actually returned could be less than

7.4 Interface: IntegrityManagement

This interface, which extends the PDController interface, allows integrity management for the Protection Domain controlled by the parent PDController interface.

The IntegrityManagement interface includes the following methods:

- `IntegrityManagement.getIntegrityLink()`
- `IntegrityManagement.getProtectedCredentials()`

7.4.1 Method: getIntegrityLink()

- **Function:** Get a certificate signed by the TCB component providing the integrity of the current Protection Domain. The signature allow links between the integrity of the TCB and the integrity of the current Protection Domain.

- **Parameters:**

Name	Type	Description
<code>nonce</code>	<code>[in] byte[]</code>	Random bytes provided by the challenger to guarantee freshness and prevent replay attacks. This nonce will be hashed by the virtualization layer using a predefined hash algorithm and included in the Integrity Report.

- **Return value:**

Name	Type	Description
	<code>IntegrityReport</code>	A report signed by the virtualization layer regarding the integrity of the Protection Domain controlled by this interface.

7.4.2 Method: getProtectedCredentials()

- **Function:** Provides Credentials to a specific Protection Domain, protected by the TCB and having the integrity requirements on the various Protection Domain enforced by the TCB.

- **Parameters:**

Name	Type	Description
<code>policy</code>	<code>[in] IntegrityPolicy</code>	The policy related to these credentials.
<code>protected_credentials</code>	<code>[in] byte[]</code>	Secret provided to the Protection Domain, protected (encrypted) by a key under the TCB control (and presumably bound to some PCRs). The result of decrypting the data should be a ProtectedCredential structure.

- **Return value:**

Name	Type	Description
	<code>byte[]</code>	The decrypted secret (if the integrity requirement is met) or an empty array if not.

7.5 Data Structures

7.5.1 Structure: ConnectionPolicy

A connection policy specifies the type of allowed connections between two Protection Domains and the direction of the connection. Connections are, at least conceptually, directed even though some implementation may not be able to actually distinguish source Protection Domains from destination Protection Domains.

Member Name	Type	Description
con_type	ConnectionType	The type of connection (if supported by the underlying hypervisor).
mode	int	Specifies the access control in terms of read/write privileges. Note: The underlying implementation might support only r/w channel and might not be able to support directed channels.

7.5.2 Structure: ConnectionIntegrityPolicy

This structure, which extends ConnectionPolicy, specifies integrity requirement on a connection policy.

Member Name	Type	Description
dest	byte[]	The integrity hash of the Protection Domain at the destination end of this connection.

7.5.3 Structure: IntegrityReport

This structure provide an external entity with a report of the integrity of a Protection Domain and link it with the integrity of the system.

Member Name	Type	Description
PDIIntegrity	byte[]	The hash produced by the Hypervisor and startup services when instantiating the Protection Domain this integrity report refers to.
pubkey	byte[]	The public part of the key used to do the signature operation.
signature	byte[]	Signature performed by the TCB over (Nonce+Protection Domain Integrity+PCRs). This could be the result of a TCG quote operation, or of a "normal" signature operation from a key bound to PCRs.

7.5.4 Structure: IntegrityPolicy

This structure allow a client to specify some requirement on the integrity of a Protection Domain (including its connections to other Protection Domains and the integrity of the underlying system). This requirement will be translated by the BMI implementation into the decision of allowing a Protection Domain to have access to its credentials or not.

Member Name	Type	Description
pdIntegrity	byte[]	The expected hash of the Protection Domain. see: IntegrityReport.PDIntegrity .
pdConnections	ConnectionIntegrityPolicy[]	List of ConnectionIntegrityPolicy from the Protection Domain with the pdIntegrity hash.

7.5.5 Structure: ProtectedCredentials

Member Name	Type	Description
secret	byte[]	A secret the Protection Domain will later try to retrieve using the IntegrityManagement interface, see: . Note: There might be a size limit on the secret depending on the size of the key used to do the encryption
policyValidation	byte[]	A hash of the IntegrityPolicy structure related to this secret. This allows the BMI implementation to link the secret for a Protection Domain with its Integrity Policy.

7.5.6 Structure: PDDescription

Description of the resource requirements for a Protection Domain.

Member Name	Type	Description
mem	int	Amount of requested Memory in bytes. Note: The underlying implementation might round it up depending on the granularity it can support. Note: If the virtualization layer can not provide the minimal amount of requested memory, the allocation should fail.
cpu	int	Number of requested CPUs. Note: Due to differences in the underlying implementation and the dynamic aspect of the resource allocations, there is no guarantee

Member Name	Type	Description
performance	<i>int</i>	that the created Protection Domain will actually have the requested number of CPUs. Indicator about the "share" of the CPUs the Protection Domain will get. Note: Due to differences in the underlying implementation and the dynamic aspect of the resource allocations, there is no guarantee that the created Protection Domain will actually have the requested performance index.

7.5.7 Structure: RuntimeCapability

This structure describes a parameter recognized by a Protection Domain, and whether its value can be changed or not. The underlying implementation of the BMI is responsible to translate the change or the parameter's value into the appropriate action (such as an event sent to the Protection Domain the parameter refers to, or to a configuration change request in the hypervisor itself). Read only capabilities are usually used to report the current state of the Protection Domain or to inform the client application calling the BMI about the supported capabilities of a particular Protection Domain.

Member Name	Type	Description
name	<i>string</i>	The name of the parameter or capability. The name follows a hierarchical naming schema similar to a file system path (like the /proc directory in standard Linux kernel). The character '/' is used as a delimiter for the hierarchy. This approach allow flexibility and extensibility without compromising compatibility. Note: The naming schema for the parameters name is defined in a separate document.
read	<i>boolean</i>	This boolean indicates if the parameter named name of the Protection Domain exposing it can be read. Non readable parameters should be writable and could be used to implement requests that translate into events. In such cases, reading the value of the corresponding parameter would be semantically wrong.
write	<i>boolean</i>	The boolean parameter indicate if the parameter name name can be

Member Name	Type	Description
		written (modified). Non writable parameters are usually used to inform the caller about the current status of the Protection Domain or its capabilities. When a parameter is modified, the BMI will translate the modification into an appropriate action.

7.6 Runtime Parameters

Runtime parameters follow a hierarchical definition similar to standard file systems. They are used to either reflect the status of a Protection Domain or to modify some of its characteristics at runtime. Each parameter is composed of a name, an optional string value (depending if the parameter can be read and/or written) and a list of child (which can be empty). A completely specified parameter name contains a list of parameter names separated by the '/' character which uniquely identifies a parameters through the hierarchy. While this naming convention simplifies the definition of the semantics of the parameters and makes it more intuitive to humans, the underlying implementation could simply rely on a flat list of completely specified parameter names.

This section describes the list of parameter that Protection Domains may expose and the semantic of each of those parameter. The implementation of the BMI will then abstract the various parameters and their modification and translate the read/write operations into the operation appropriate for a specific implementation of the TVL. Note however that the caller of the BMI would have to check for the actual supported parameters using the Controller interface, and should not expect the existence of all the parameters for every Protection Domain. The caller should also check if a particular parameter can be read and/or written as this capability might also depends on the implementation of the TVL and the type of the Protection Domain associated to this parameter.

Finally, this approach of naming parameters has been chosen for its flexibility and extensibility, so as a consequence, the currently provided list is likely to be modified and increase over the duration of the project.

Parameter Name (path)	Description
/system/memory/max	Maximal amount of memory the Protection Domain is allowed to use.
/system/memory/min	Minimal amount of memory the Protection Domain will ever be allocated.
/system/memory/current	The amount of memory currently allocated to the Protection Domain
/system/cpu/max	Maximal number of CPU (or Virtual CPUs) the Protection Domain is allowed to

Parameter Name (path)	Description
	allocate.
/system/cpu/min	Minimal number of CPU (or Virtual CPUs) the Protection Domain will ever be allocated.
/system/cpu/current	The number of CPU (or Virtual CPUs) the Protection Domain current has allocated.
/capability/pause	If this parameter exists, it indicates that the Protection Domain can be paused. The value of this parameter indicates the value to send to the requestStatusChange() method of the PDController interface.
/capability/restart	If this parameter exists, it indicates that the Protection Domain can be restarted. This would be translated into a reboot operation for a virtualized operating system for instance. The value of this parameter indicates the value to send to the requestStatusChange() method of the PDController interface.
/capability/suspend	If this parameter exists, it indicates that the Protection Domain can be suspended. The value of this parameter indicates the value to send to the requestStatusChange() method of the PDController interface. The needed state will be saved by the TVL in a transparent manner. <u>Note:</u> The existence of this parameter implies the existence of "/capability/resume"
/capability/resume	If this parameter exists, it indicates that the Protection Domain can be resumed (following a suspend operation). The value of this parameter indicates the value to send to the requestStatusChange() method of the PDController interface. The needed state will be loaded by the TVL in a transparent manner. <u>Note:</u> The existence of this parameter implies the existence of "/capability/resume"

Parameter Name (path)	Description
/interfaces/#name (where #name is specific to each type of Protection Domain)	This parameter is a place holder for definition of the type of services a Protection Domain can provide. The remaining of the naming of the hierarchy will be defined by each type of services being implemented by WP4 and WP5.

This table describes the minimal commonality between all the Protection Domain, but more complex Protection Domain will define extra entry in this hierarchy. For instance, a virtualized operating system may provide a “/device” entry to list the devices it is connected to and the corresponding Protection Domain local ID of the device driver.

8 List of Abbreviations / Glossary

BBB	BIOS boot block
BMI	Basic Management Interface
bmodfs	L4 specific: implements file provider with generic_fprov interface
CIM	Common Information Model
CRTM	Core Root of Trusted Measurement
generic_fprov	L4 specific: file provider interface definition
iSCSI	Internet SCSI
L4 Con	L4 specific: Lightweight graphical console
MAC	Mandatory Access Control
MBR	Master Boot Record
NBD	Network Block Device
NFS	Network File System
PCI	Peripheral Component Interface
PCR	Platform Configuration Register
RTM	Root of Trusted Measurement
run	Component controlling the loader (requires L4 Con)
SCSI	Small Computer System Interface
TCG	Trusted Computing Group
tftp	L4 specific: component implementing generic_fprov, uses tftp server over network
TPM	Trusted Platform Module
TVL	Trusted Virtualization Layer
VBD	Virtual Block Device
VCPU	Virtual CPU